

Just like in econometrics, the most basic “learning algorithm” is OLS

Just like in econometrics, the most basic “learning algorithm” is OLS

Or, if Y is categorical, logistic regression

Just like in econometrics, the most basic “learning algorithm” is OLS

Or, if Y is categorical, logistic regression

But, we know there are many other ways to estimate models

Just like in econometrics, the most basic “learning algorithm” is OLS

Or, if Y is categorical, logistic regression

But, we know there are many other ways to estimate models

e.g. non-parametric, semi-parametric, Bayesian, ...

- Supervised learning: we predict Y from a set of X 's

- Supervised learning: we predict Y from a set of X 's
- Unsupervised learning: we try to group observations by their X 's

- Supervised learning: we predict Y from a set of X 's
- Unsupervised learning: we try to group observations by their X 's
- Most of econometrics is about supervised learning (i.e. estimate $\hat{\beta}$)

- Supervised learning: we predict Y from a set of X 's
- Unsupervised learning: we try to group observations by their X 's
- Most of econometrics is about supervised learning (i.e. estimate $\hat{\beta}$)
- But there are some elements of unsupervised learning

- Supervised learning: we predict Y from a set of X 's
- Unsupervised learning: we try to group observations by their X 's
- Most of econometrics is about supervised learning (i.e. estimate $\hat{\beta}$)
- But there are some elements of unsupervised learning
 - Particularly with regards to detecting unobserved heterogeneity types

- Supervised learning: we predict Y from a set of X 's
- Unsupervised learning: we try to group observations by their X 's
- Most of econometrics is about supervised learning (i.e. estimate $\hat{\beta}$)
- But there are some elements of unsupervised learning
 - Particularly with regards to detecting unobserved heterogeneity types
 - e.g. factor analysis (detect types based on a set of measurements)

- Supervised learning: we predict Y from a set of X 's
- Unsupervised learning: we try to group observations by their X 's
- Most of econometrics is about supervised learning (i.e. estimate $\hat{\beta}$)
- But there are some elements of unsupervised learning
 - Particularly with regards to detecting unobserved heterogeneity types
 - e.g. factor analysis (detect types based on a set of measurements)
 - e.g. the EM algorithm (detect types based on serial correlation of residuals)

Examples of supervised learning algorithms:

Examples of supervised learning algorithms:

- Tree models

Examples of supervised learning algorithms:

- Tree models
 - Basically a fully non-parametric bin estimator

Examples of supervised learning algorithms:

- Tree models
 - Basically a fully non-parametric bin estimator
 - Can generalize to “forests” that average over many “trees”

Examples of supervised learning algorithms:

- Tree models
 - Basically a fully non-parametric bin estimator
 - Can generalize to “forests” that average over many “trees”
- Neural networks

Examples of supervised learning algorithms:

- Tree models
 - Basically a fully non-parametric bin estimator
 - Can generalize to “forests” that average over many “trees”
- Neural networks
 - Model the human brain’s system of axons and dendrites

Examples of supervised learning algorithms:

- Tree models
 - Basically a fully non-parametric bin estimator
 - Can generalize to “forests” that average over many “trees”
- Neural networks
 - Model the human brain’s system of axons and dendrites
 - “Input layer” is the X ’s, “Output layer” is Y

Examples of supervised learning algorithms:

- Tree models
 - Basically a fully non-parametric bin estimator
 - Can generalize to “forests” that average over many “trees”
- Neural networks
 - Model the human brain’s system of axons and dendrites
 - “Input layer” is the X ’s, “Output layer” is Y
 - “Hidden layers” nonlinearly map the input and output layers

Examples of supervised learning algorithms:

- Tree models
 - Basically a fully non-parametric bin estimator
 - Can generalize to “forests” that average over many “trees”
- Neural networks
 - Model the human brain’s system of axons and dendrites
 - “Input layer” is the X ’s, “Output layer” is Y
 - “Hidden layers” nonlinearly map the input and output layers
 - Mapping ends up looking like a logit of logit of logits

Examples of supervised learning algorithms:

- Tree models
 - Basically a fully non-parametric bin estimator
 - Can generalize to “forests” that average over many “trees”
- Neural networks
 - Model the human brain’s system of axons and dendrites
 - “Input layer” is the X ’s, “Output layer” is Y
 - “Hidden layers” nonlinearly map the input and output layers
 - Mapping ends up looking like a logit of logit of logits
 - Linear regression is a special case

- Bayesian models

- Bayesian models
 - Bayes' rule can be thought of as a learning algorithm

- Bayesian models
 - Bayes' rule can be thought of as a learning algorithm
 - Use it to update one's prior

- Bayesian models
 - Bayes' rule can be thought of as a learning algorithm
 - Use it to update one's prior
- Support Vector Machine (SVM)

- Bayesian models
 - Bayes' rule can be thought of as a learning algorithm
 - Use it to update one's prior
- Support Vector Machine (SVM)
 - Originally developed for classification

- Bayesian models
 - Bayes' rule can be thought of as a learning algorithm
 - Use it to update one's prior
- Support Vector Machine (SVM)
 - Originally developed for classification
 - Tries to separate 0s and 1s by as large of a margin as possible

- Bayesian models
 - Bayes' rule can be thought of as a learning algorithm
 - Use it to update one's prior
- Support Vector Machine (SVM)
 - Originally developed for classification
 - Tries to separate 0s and 1s by as large of a margin as possible
 - Based on representing examples as points in space

- Bayesian models
 - Bayes' rule can be thought of as a learning algorithm
 - Use it to update one's prior
- Support Vector Machine (SVM)
 - Originally developed for classification
 - Tries to separate 0s and 1s by as large of a margin as possible
 - Based on representing examples as points in space
 - Generalization of the maximal margin classifier

We covered the EM algorithm and PCA in previous videos

We covered the EM algorithm and PCA in previous videos

k-means clustering

We covered the EM algorithm and PCA in previous videos

k-means clustering

- Attempts to group observations together based on the X 's

We covered the EM algorithm and PCA in previous videos

k-means clustering

- Attempts to group observations together based on the X 's
- Choose cluster labels to minimize difference in X 's among labeled observations

We covered the EM algorithm and PCA in previous videos

k-means clustering

- Attempts to group observations together based on the X 's
- Choose cluster labels to minimize difference in X 's among labeled observations

$$\min_{C_1, \dots, C_K} \sum_{k=1}^K \frac{1}{N_k} \sum_{i \in C_k} \sum_{\ell=1}^L (x_{i\ell} - \bar{x}_{\ell j})^2$$

We covered the EM algorithm and PCA in previous videos

k -means clustering

- Attempts to group observations together based on the X 's
- Choose cluster labels to minimize difference in X 's among labeled observations

$$\min_{C_1, \dots, C_K} \sum_{k=1}^K \frac{1}{N_k} \sum_{i \in C_k} \sum_{\ell=1}^L (x_{i\ell} - \bar{x}_{\ell j})^2$$

N_k is the number of observations in cluster k , L is number of X 's

We covered the EM algorithm and PCA in previous videos

k -means clustering

- Attempts to group observations together based on the X 's
- Choose cluster labels to minimize difference in X 's among labeled observations

$$\min_{C_1, \dots, C_K} \sum_{k=1}^K \frac{1}{N_k} \sum_{i \in C_k} \sum_{\ell=1}^L (x_{i\ell} - \bar{x}_{\ell j})^2$$

N_k is the number of observations in cluster k , L is number of X 's

Can choose other metrics besides Euclidean distance

- Active learning: algorithm chooses the next example it wants a label for

- Active learning: algorithm chooses the next example it wants a label for
- Balances “exploration” and “exploitation”

- Active learning: algorithm chooses the next example it wants a label for
- Balances “exploration” and “exploitation”
- Two common examples of active learning:

- Active learning: algorithm chooses the next example it wants a label for
- Balances “exploration” and “exploitation”
- Two common examples of active learning:
 1. Reinforcement learning powers the world-beating chess engines

- Active learning: algorithm chooses the next example it wants a label for
- Balances “exploration” and “exploitation”
- Two common examples of active learning:
 1. Reinforcement learning powers the world-beating chess engines
 - These algorithms use dynamic programming methods

- Active learning: algorithm chooses the next example it wants a label for
- Balances “exploration” and “exploitation”
- Two common examples of active learning:
 1. Reinforcement learning powers the world-beating chess engines
 - These algorithms use dynamic programming methods
 - Use Conditional Choice Probabilities for computational gains

- Active learning: algorithm chooses the next example it wants a label for
- Balances “exploration” and “exploitation”
- Two common examples of active learning:
 1. Reinforcement learning powers the world-beating chess engines
 - These algorithms use dynamic programming methods
 - Use Conditional Choice Probabilities for computational gains
 2. Recommender systems power social networks, streaming services, etc.